

# What are Production systems in Kubernetes?

In a production environment, Kubernetes (K8s) is used to manage containerized applications at scale. Several solutions and strategies are commonly employed to ensure reliability, high availability, scalability, and ease of management for production-grade Kubernetes clusters.



## Managing Kubernetes clusters in production systems

- DevOps engineers are responsible for creating, configuring and deleting Kubernetes clusters in organizations
- It is important to understand Kubernetes distribution and popular options
- There is a difference between development and production environments
- Interview questions may focus on managing Kubernetes clusters in production
- Paid Kubernetes distributions like Red Hat or Amazon Linux ensure timely security updates, similar to paid Linux distributions

## Production systems and tools often used with Kubernetes:

Production-grade Kubernetes clusters require a robust ecosystem of tools and solutions to ensure reliability, high availability, scalability, and efficient management. Below are some of the production systems and commonly used tools in conjunction with Kubernetes



### 1. Kubernetes Distributions:

- **Amazon Elastic Kubernetes Service (EKS):** Managed Kubernetes service by AWS.
- **Google Kubernetes Engine (GKE):** Managed Kubernetes service by Google Cloud.
- **Azure Kubernetes Service (AKS):** Managed Kubernetes service by Microsoft Azure.

- **OpenShift:** An enterprise Kubernetes platform by Red Hat.
- **Rancher:** An open-source Kubernetes management platform.

## 2. Container Runtimes:

- **Docker:** Historically, Docker was the most commonly used container runtime, but Kubernetes can work with various runtimes.
- **containerd:** An industry-standard core container runtime used by Kubernetes.
- **CRI-O:** An OCI-compliant container runtime optimized for Kubernetes.

## 3. Cluster Networking:

- **Flannel:** A simple and easy-to-set-up overlay network solution for Kubernetes.
- **Calico:** A networking and network security solution for Kubernetes.
- **Weave:** A lightweight network solution for Kubernetes that simplifies network management.

## 4. Ingress Controllers:

- **NGINX Ingress Controller:** Manages ingress resources and routes traffic to services.
- **Traefik:** A modern reverse proxy and load balancer.
- **HAProxy:** A high-performance TCP/HTTP load balancer.

# what is KOPS ?

Kops, short for **Kubernetes Operations** is an open-source command-line tool that facilitates the provisioning, deployment, and management of production-ready Kubernetes clusters on cloud infrastructure. It is particularly well-suited for creating and managing Kubernetes clusters on cloud providers like Amazon Web Services (AWS), Google Cloud Platform (GCP), and Microsoft Azure, but it can also be used on other platforms.



Key features and capabilities of Kops include:

1. **Cluster Lifecycle Management:** Kops simplifies the entire lifecycle of a Kubernetes cluster, including cluster creation, scaling, updates, and deletion. It allows you to manage the cluster like any other infrastructure as code (IAC).
2. **Infrastructure Provisioning:** It automates the provisioning of the underlying infrastructure components necessary for a Kubernetes cluster, such as virtual machines, networking, and storage. Kops can create

and configure these resources on various cloud platforms.

3. **Cluster Customization:** Kops provides fine-grained control over the configuration of your Kubernetes cluster. You can specify details like instance types, network settings, and add-ons to tailor the cluster to your specific requirements.
4. **High Availability:** Kops supports the creation of highly available Kubernetes clusters by distributing control plane components across multiple availability zones to ensure fault tolerance.
5. **Upgrade Management:** It simplifies the process of upgrading your Kubernetes cluster to newer versions, ensuring that your clusters are up to date with the latest security patches and features.
6. **Validation and Verification:** Kops offers commands to validate your cluster's configuration and state, helping you ensure that your cluster is correctly configured and functioning as expected.
7. **Integration with DNS:** Kops can automatically configure DNS entries and integration with cloud DNS services to make it easier to access services running on your Kubernetes cluster.
8. **Backup and Restore:** It provides features for creating backups of your cluster's state, allowing you to restore your cluster in case of failures or mistakes.
9. **Community and Open Source:** Kops is an open-source project with an active community of contributors and users. It benefits from regular updates and improvements driven by the Kubernetes community.

## Installation Kubernetes on Amazon EC2 instances using KOPS.



Detail to install Kubernetes on Amazon EC2 instances using KOPS.

### Prerequisites

Before you start, ensure you have the required prerequisites:

1. **Python 3:** Python is a programming language, and Python 3 is a specific version. You can check if Python 3 is installed by running `python3 --version`. If it's not installed, you can typically install it using your operating system's package manager.
2. **AWS CLI (Command Line Interface):** The AWS CLI is a command-line tool that allows you to interact with AWS services from your terminal. You can install it by following the instructions in the [official AWS CLI documentation \(https://docs.aws.amazon.com/cli/latest/userguide/install-cliv2.html\)](https://docs.aws.amazon.com/cli/latest/userguide/install-cliv2.html).
3. **kubectl:** Kubectl is the command-line tool for interacting with Kubernetes clusters. You can install it using your package manager or by downloading it directly from the [Kubernetes releases page](#)

[\(https://kubernetes.io/docs/tasks/tools/install-kubectl/\)](https://kubernetes.io/docs/tasks/tools/install-kubectl/).

## Installation of Dependencies

These commands install the necessary software dependencies on your EC2 instance or laptop:

```
# Add the Kubernetes APT repository key
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
```

This command retrieves a GPG key used to verify the authenticity of the Kubernetes APT repository.

```
# Add the Kubernetes APT repository
echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/sources.list.d/kubernetes.list
```

Here, you're adding the Kubernetes APT repository to your system's list of package sources.

```
# Update package lists and install required packages
sudo apt-get update
sudo apt-get install -y python3-pip apt-transport-https kubectl
```

These commands perform the following tasks:

- Update the package lists to fetch the latest information about available packages.
- Install Python 3, which is necessary for some tools and scripts.
- Install the `apt-transport-https` package, which enables the use of HTTPS for downloading packages over the network.
- Install `kubectl`, the Kubernetes command-line tool, which is used to interact with your Kubernetes cluster.

```
# Upgrade AWS CLI to the latest version
pip3 install awscli --upgrade
```

This command upgrades the AWS CLI to the latest version to ensure you have access to the latest features and bug fixes.

```
# Add AWS CLI to the PATH
export PATH="$PATH:/home/ubuntu/.local/bin/"
```

This line adds the AWS CLI executable to your system's PATH so that you can use it from any location in your terminal.



## Installation of KOPS

KOPS (Kubernetes Operations) is a tool for managing Kubernetes clusters on AWS. These commands install KOPS on your system:

```
# Download the latest version of KOPS for Linux
curl -LO https://github.com/kubernetes/kops/releases/download/$(curl -s http
s://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut
-d '"' -f 4)/kops-linux-amd64
```

Here's what each part does:

- `curl -s https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut -d '"' -f 4` fetches the latest release tag for KOPS from GitHub.
- `curl -LO` downloads the KOPS binary for Linux.

```
# Make KOPS executable
chmod +x kops-linux-amd64
```

This command makes the KOPS binary executable so you can run it as a command.

```
# Move KOPS to a location in the PATH
sudo mv kops-linux-amd64 /usr/local/bin/kops
```

This line moves the KOPS binary to `/usr/local/bin`, which is in the system's PATH, so that you can run it from any location in your terminal.

### Add user

1 2 3 4 5

#### Set permissions

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

AdministratorAccess

Showing 4 results

	Policy name	Type	Used as
<input checked="" type="checkbox"/>	AdministratorAccess	Job function	Permissions policy (8)
<input type="checkbox"/>	AdministratorAccess-Amplify	AWS managed	None
<input type="checkbox"/>	AdministratorAccess-AWSElasticBeanstalk	AWS managed	None
<input type="checkbox"/>	AWSAuditManagerAdministratorAccess	AWS managed	None

## AWS IAM Permissions

AWS IAM (Identity and Access Management) is used to manage access to AWS services. Ensure that the IAM user you are using has the following permissions:

1. **AmazonEC2FullAccess**: This permission allows the user to create and manage EC2 instances, which are used for hosting your Kubernetes cluster nodes.
2. **AmazonS3FullAccess**: S3 (Simple Storage Service) is used for storing KOPS configuration and cluster state. This permission is required for KOPS to manage these resources.
3. **IAMFullAccess**: This permission allows the user to manage IAM roles and policies, which are necessary for AWS services to interact securely with each other.
4. **AmazonVPCFullAccess**: Amazon VPC (Virtual Private Cloud) is used for networking and security configuration of your EC2 instances. This permission is required to set up networking for your Kubernetes cluster.

## AWS CLI Configuration

Before you can use the AWS CLI, you need to configure it with your AWS credentials. Run the following command:

```
aws configure
```

This command will prompt you to enter your AWS Access Key ID, AWS Secret Access Key, default region name, and default output format. These credentials will be used by the AWS CLI to authenticate your requests to AWS services.

## Kubernetes Cluster Installation

Now, let's proceed with installing a Kubernetes cluster on AWS using KOPS.

### Create an S3 Bucket

Before creating the cluster, you need to create an S3 bucket that KOPS will use to store cluster configuration and state. Run this command:

```
aws s3api create-bucket --bucket kops-Prudhvi-storage --region us-east-1
```

This command creates an S3 bucket named `kops-abhi-storage` in the `us-east-1` region. You can replace these values with your preferred bucket name and region.

### Create the Kubernetes Cluster

Use the `kops create cluster` command to create a Kubernetes cluster. This command specifies various parameters for the cluster, such as its name, state storage location, AWS availability zone, node configurations, and more. Here's an example command:

```
kops create cluster --name=demok8scluster.k8s.local --state=s3://kops-Prudhvi-s  
torage --zones=us-east-1a --node-count=1 --node-size=t2.micro --master-size=t2.  
micro --master-volume-size=8 --node-volume-size=8
```

Let's break down the important options:

- `--name` : The name of your Kubernetes cluster.
- `--state` : The S3 bucket location where KOPS stores the cluster's state and configuration.
- `--zones` : The AWS availability zones where the cluster's nodes will be deployed.
- `--node-count` : The number of worker nodes (EC2 instances) in your cluster.
- `--node-size` and `--master-size` : The EC2 instance types for your nodes and master node.
- `--master-volume-size` and `--node-volume-size` : The size of the root volumes attached to the master and node instances.

### Edit the Cluster Configuration

Before proceeding further, you should edit the cluster configuration to review and adjust the settings according to your needs. Use the `kops edit cluster` command, as follows:

```
kops edit cluster myfirstcluster.k8s.local
```

Replace `myfirstcluster.k8s.local` with the actual name of your cluster.

This

---

## Kubernetes on Local Development Environments

Running Kubernetes (K8s) on local development environments is essential for development and testing before deploying applications to production clusters. Several tools and approaches are available to set up Kubernetes on your local machine for development purposes.



### 1. Minikube:

- **Description:** Minikube is a lightweight and easy-to-install tool that allows you to run a single-node Kubernetes cluster on your local machine. It's designed for local development and testing.
- **Features:**

- Provides a complete Kubernetes environment.
  - Supports various container runtimes, including Docker and containerd.
  - Easily switch between different Kubernetes versions.
  - Integration with kubectl for managing the local cluster.
- **Installation:** Install Minikube according to your operating system by following the official documentation: <https://minikube.sigs.k8s.io/docs/start/> (<https://minikube.sigs.k8s.io/docs/start/>)



## 2. Kind (Kubernetes in Docker):

- **Description:** KinD is a tool that enables you to create Kubernetes clusters using Docker containers as nodes. It's suitable for testing and development.
- **Features:**
  - Lightweight and fast cluster creation.
  - Ideal for creating multi-node clusters on a single machine.
  - Easy integration with existing Docker workflows.
  - Good support for testing Kubernetes features.
- **Installation:** Install Kind by following the instructions for your platform: <https://kind.sigs.k8s.io/docs/user/quick-start/> (<https://kind.sigs.k8s.io/docs/user/quick-start/>)



## 3. K3d (K3s in Docker):

- **Description:** K3d is a lightweight wrapper for K3s, a lightweight Kubernetes distribution. It allows you to run K3s clusters in Docker containers.
- **Features:**
  - Quick and lightweight cluster creation.
  - Minimal resource consumption.
  - Suitable for local testing and development.
- **Installation:** Install K3d by following the instructions in the GitHub repository: <https://github.com/rancher/k3d> (<https://github.com/rancher/k3d>)





# MicroK8s

## 4. MicroK8s:

- **Description:** MicroK8s is a Kubernetes distribution designed for easy installation and operation on Linux systems. It provides a full-featured Kubernetes environment.
- **Features:**
  - Lightweight and fast setup.
  - Supports multi-node clusters on a single machine.
  - Includes add-ons like Istio, Helm, and more.
  - Suitable for both local development and edge computing.
- **Installation:** Install MicroK8s according to the documentation for your Linux distribution: <https://microk8s.io/docs> (<https://microk8s.io/docs>).



## 5. Kubeadm (Advanced Option):

- **Description:** Kubeadm is a tool for bootstrapping Kubernetes clusters. While it's typically used for setting up production clusters, it can also be used to create a single-node cluster on your local machine.
- **Features:**
  - Full control over cluster configuration.
  - Suitable for advanced users and those interested in understanding Kubernetes internals.
- **Installation:** Set up a single-node Kubernetes cluster using Kubeadm by following the Kubernetes documentation: <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/> (<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>).

## Example : Installation of Minikube

Installing Minikube is a straightforward process, and it allows you to run a single-node Kubernetes cluster on your local machine. Below are detailed installation steps for Minikube on Linux (Ubuntu-based systems). Be sure to check the [official Minikube documentation](https://minikube.sigs.k8s.io/docs/start/) (<https://minikube.sigs.k8s.io/docs/start/>) for the most up-to-date instructions for your specific OS.

### 1. Prerequisites:

Before you begin, ensure that you have the following prerequisites:

- A working Linux environment (these instructions are for Ubuntu-based systems).
- A hypervisor installed (e.g., VirtualBox, KVM, or Docker, although Docker support may be deprecated).
- Kubectl (Kubernetes command-line tool) installed. If not installed, you can install it with the following command:

```
sudo snap install kubectl --classic
```

## 2. Install VirtualBox (if not already installed):

If you don't have VirtualBox installed, you can install it using the package manager:

```
sudo apt-get update  
sudo apt-get install virtualbox virtualbox-ext-pack
```

Follow the prompts to complete the installation, including accepting the license agreement for the VirtualBox Extension Pack.

## 3. Install Minikube:

You can install Minikube using the package manager Snap. First, update your package list:

```
sudo apt-get update
```

Then, install Minikube:

```
sudo snap install minikube --classic
```

## 4. Start Minikube:

Now that Minikube is installed, you can start a Minikube cluster using the `minikube start` command. You can specify additional options if needed, such as the Kubernetes version:

```
minikube start --kubernetes-version=<desired-version>
```

Replace `<desired-version>` with the Kubernetes version you want to use (e.g., `v1.22.0` ).

Minikube will download the specified Kubernetes version and start a virtual machine to host the cluster.

## 5. Check Minikube Status:

After a successful start, you can check the status of your Minikube cluster using the following command:

```
minikube status
```

This command will display information about the cluster's status, such as the Kubernetes version and whether it's running.

## 6. Interact with Minikube:

You can interact with your Minikube cluster using `kubectl`. The Minikube binary automatically configures `kubectl` to use the Minikube cluster context. To check the cluster's nodes, use:

```
kubectl get nodes
```

This should display your Minikube node as "Ready."

## 7. Stop and Delete Minikube Cluster:

When you're done working with your Minikube cluster, you can stop and delete it using the following commands:

```
minikube stop  
minikube delete
```

This will stop the cluster and remove the associated virtual machine.

That's it! You've successfully installed Minikube on your Ubuntu-based Linux system and started a local Kubernetes cluster. You can now use this cluster for development and testing purposes.

-- Prudhvi Vardhan (LinkedIn)